BERT-Enhanced Automated Pattern Recognition in Software Repositories Using Natural Language Processing and Machine Learning

Sai Krishna Cholleti*1, Bandi Nikhil*2, Pappula Amulya*3

[1]Artificial Intelligence & Machine Learning, CMR College of Engineering& Technology, Hyderabad, Telangana, India
saikrishna.cholleticeh@gmail.com [1]

[2]Artificial Intelligence & Machine Learning, CMR College of Engineering& Technology, Hyderabad, Telangana, India
nikhilbandii2205@gmail.com [2]

[3] Artificial Intelligence & Machine Learning, CMR College of Engineering& Technology, Hyderabad, Telangana, India
pappulareddy8@gmail.com [3]

ABSTRACT

The sheer volume of textual data created by software repositories (e.g., bug reports, developer comments, commit messages) is becoming increasingly challenging to manually analyze due to the increasing size and complexity of software projects. As such, existing text-mining and traditional machine-learning approaches to analyzing these datasets rely primarily on the frequency of keywords and do not capture the contextual meaning behind technical descriptions. In response to this challenge, this research presents an automated pattern-recognition approach using a Bidirectional Encoder Representations from Transformers (BERT)-based model to leverage advances in Natural Language Processing (NLP), and Machine Learning (ML) to analyze and effectively interpret software-repository data. Using contextual embeddings from the BERT model this proposed method can be used to automatically classify bugs, find common error patterns in the code base, predict severity of errors based on their frequency of occurrence, and even find duplicate bug reports by analyzing similarities between them. Additionally, the system includes confidence-scores, root cause-hints, and bug-trend insight at the repository level through a Web-Based User Interface. Experimental results show that the proposed framework enhances the developers' ability to semantically understand their source-code, reduces the manually required work, increases the speed and accuracy of software-maintenance activities. Therefore, the results of this research illustrate the practicality of utilizing transformer-based models to automate the analysis of software-repositories, and enhance decision-making by software-development teams.

Keywords: BERT(Bidirectional Encoder Representations from Transformers), Software Repository Mining, Bug Severity Prediction, Duplicate Bug Detection, Text Clustering, NLP-Based Software Maintenance

I. INTRODUCTION

Developers encounter massive textual data during the creation of software through repositories such as Bugzilla, GitHub, and Jira. These repositories hold messages, code comments, and bug reports, which reflect the continuous evolution and quality of the system. Due to the complexity of projects, manual analysis is often impractical and inefficient. The lack of prioritization and bug identification leads to increased costs and decreased software quality.

The traditional methods for software repository analysis heavily depend on keyword- based representation (like TF-IDF) and classical machine learning classifiers. These approaches while offering simple automation lack semantical relation and context in technical text. This has led to issues being reported more than once, employees or users not finding relevant reports, and incorrect assignment of priority. Popular natural language processing transformers dramatically elevate the level of understanding of text in context. BERT supports bidirectional context modeling, and is more powerful in comprehending the semantic meanings of technical descriptions than traditional models.

As a matter of fact, BERT proves to be a relatively cool tool to extract information about software repositories since it considers the context and the intent of the developers to express. It fits very well in the process of trying to comprehend everything that goes on in and around bugs, commits and discussion cards.

Despite all those updates, many of the available systems continue to be more narrow-purpose such as bug sorting or dupe detection, and they tend to provide you with a single view of the entire repository. It implies that there is a lot of work ahead on the development of a unified, automated system, which will be able to draw in

semantic analysis, severity prediction, similarity identification, and other, larger analytics.

1 builds the gap that researchers are attempting to bridge by generating a single and automated pipeline capable of integrating all those parts. As such, this paper attempts to address this gap by introducing a software-repositories pattern-recognition framework that is powered by BERT. The process takes raw text and converts the text into BERT embeddings, clusters patterns, predicts the severity of bugs depending on their frequency based on frequency, and identifies duplicates based on a similarity check. Transparency is also provided by the system with the addition of confidence scores, root-cause hints and repository-scale trend insights. Interpreted primarily, it is evident that the authors are seeking to present a scalable solution that ensures the software maintenance is easier and the clarity of the decision-making process.

II. RELATED WORK

Studies related to the analysis of software repositories have been popular recently because of the tremendous amount of text created during the software development process. A number of valuable pieces of information about the quality of a software system and maintenance issues can be obtained from software repositories through the bug reports, commit messages and developer discussion contained within them. Early research mainly concentrated on automatically classifying bugs, automatically classifying bug reports, and automatically detecting duplicates in bug reports using traditional text mining and machine learning techniques. While these early studies provided the basis for the use of automation in analyzing software repositories, they had significant disadvantages in dealing with the sematic complexities as well as the context variations found in technical text.

Traditional Text-Based Approaches

The initial efforts employed Bag -of-Words and TF - IDF with Naïve Bayes, SVMs, and decision trees as the classifier in order to classify and filter bugs. They did a good job when the bug reports were structured in a similar fashion but were very dependent on the frequency of words and lost concept and order of words altogether. This meant that semantically similar bug reports using different words were usually misclassified and duplicate detection was a random success. Rule- based or heuristic duplicate detection with cosine similarity and constant thresholds was also attempted by many of the researchers. Then those ideas reduced the redundancy a little, but were greatly reduced in case of paraphrasing or tiny bugs. These were rapid but had less depth and scalability that is required in large and diverse repositories.

No.	Title	Problem	Solution	Limitation	Methodology
1	Bug Report Assignment	Manual triage inefficiency	Automated assignment system	Limited semantic understanding	TF-IDF + SVM
2	Developer Recommendation	Incorrect bug assignee	Expertise prediction model	Requires historical data	Supervised ML
3	Duplicate Detection	Redundant bug reports	Similarity-based filtering	Keyword dependency	Cosine similarity
4	NLP Defect Detection	Duplicate defect identification	Text similarity model	Shallow lexical features	TF-IDF + NLP
5	Topic-Based Triage	Large-scale bug triage	LDA-based clustering	Low fine-grained accuracy	Topic modelling
6	Severity Prediction	Incorrect severity labelling	ML severity classifier	Requires labelled dataset	Supervised ML
7	Bug Localization (DL)	Inaccurate fault localization	Deep neural localization	High computational cost	CNN-based DL
8	Fault Prediction	Software defect forecasting	Change history analysis	Ignores semantic context	Statistical modelling
9	Neural Severity Model	Severity estimation errors	Neural classification	Limited interpretability	Deep learning
10	Transformer Classification	Lack of contextual analysis	BERT-based embedding	Task-specific scope	Transformer model

Clustering and Topic Modeling Methods

Here is the outline of what has been accomplished so far by mapping and analyzing the data: The entire data set accounts the connection between the present-day reality and the past events of the 20th century of the United States. To overcome the constraints of the supervised learning folks proceeded to unsupervised clustering of data and topic models like the K -Means and LDA. The concept here was that the bug reports could be sorted using textual similarity and latent topics such that recurring patterns of errors could be identified without having to label data. Clustering assisted in viewing high-level trends and venture data on repositories. Nonetheless, these approaches were pegged on superficial textual characteristics and were not capable of separating semantically similar, nevertheless lexically dissimilar bug reports. The topic models would tend to generate very wide and difficult-to-contextualize collections, and were not very effective at fine-grained tasks such as severity prediction or finding accurate duplicates. That is why they were not very popular in maintenance configurations of the real world.

Deep Learning and Neural Network Approaches

The final wave involved a deep learning architecture such as CNNs and RNNs being used in classifying bugs and predicting the severity. They have enhanced it by training distributed text representations and local context. They were more specific and generalizable in comparison with older approaches. The downside? Text in CNNs and RNNs is processed sequentially, meaning that they have trouble with long-range dependence as in long and technical bug reports. That is where Transformer-based models, especially BERT, tilted the playing field by creating the opportunity to understand

the context in both directions. Some papers demonstrated BERT to perform better than conventional and neural ones on Bug classification and Bug localization.Despite all that, most BERT based systems continue to keep doing one job at a time and lack the integrated analytics you would like, such as repo-level trend analysis, scoring, or automated root-cause confidence interpretation.

III. METHODOLOGY

Methodology: Structure and Content

This paper established a coherent, stepwise process of developing a clarifiable heart -murmur screening program. It was intended to use the deep-learning acoustic model with language-based clinical interpretation under a single roof. The design focused on simplifying and explainable and screening-oriented advice as opposed to diagnostic claims through the design of the architecture.

3.1 Proposed Framework Overview

The proposed system employed a multi-layered pipeline, which converted heart-sound samples to outputs that can be read like a screening report. It consisted of five logical stages, which are audio standardization, spectral representation, deep feature learning, probability-based risk stratification, and language-directed explanation generation. The heart sounds have been reduced to a regular time -frequency format to minimize variance. Then those representations were then processed by a Convolutional Recurrent Neural Network (CRNN) with attention to estimate the likelihood of a murmur. Discrete screening risk levels were mapped to the probability. Lastly, a Large Language Model (LLM) generated systematic clinical explanations in accordance to the level of confidence of the model and the context of the patient. The process was made transparent in a way that the main predictions were not affected by the separation of inference and explanation.

3.2 System Architecture

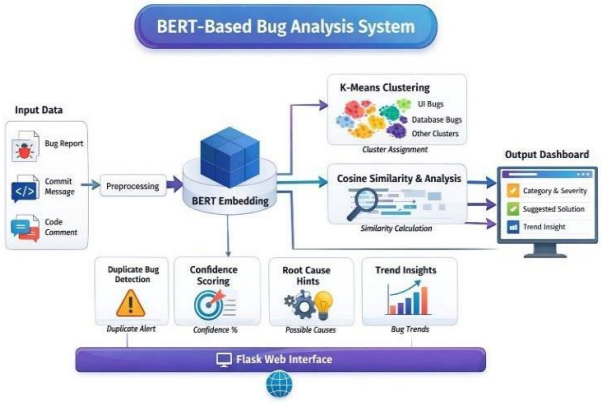

This system was constructed to be a pipeline processing similar to that of a real world design. The raw heart-sound signal was converted into Mel spectrogram by the first block that retained frequency and time resolution required to detect the murmurs.

Those tensors were inputted into convolutional layers which extracted local frequency-time abnormal blood flow features. The resulting feature maps were re-modeled and fed to bidirectional recurrent layers which model temporal dependencies between heart cycles and learn forward and backward events. That assisted in processing rhythm change and different signal lengths. A attention system was superimposed on top of the recurring outputs to weigh the likely-to-have frames of the reports that had details of the murmer. Attention-weighted representation was added up and passed to fully connected layers yielding one murmur probability by using sigmoid activation. This figure is the degree to which the model is certain of the presence of a murmur.

3.3 Processing Flow

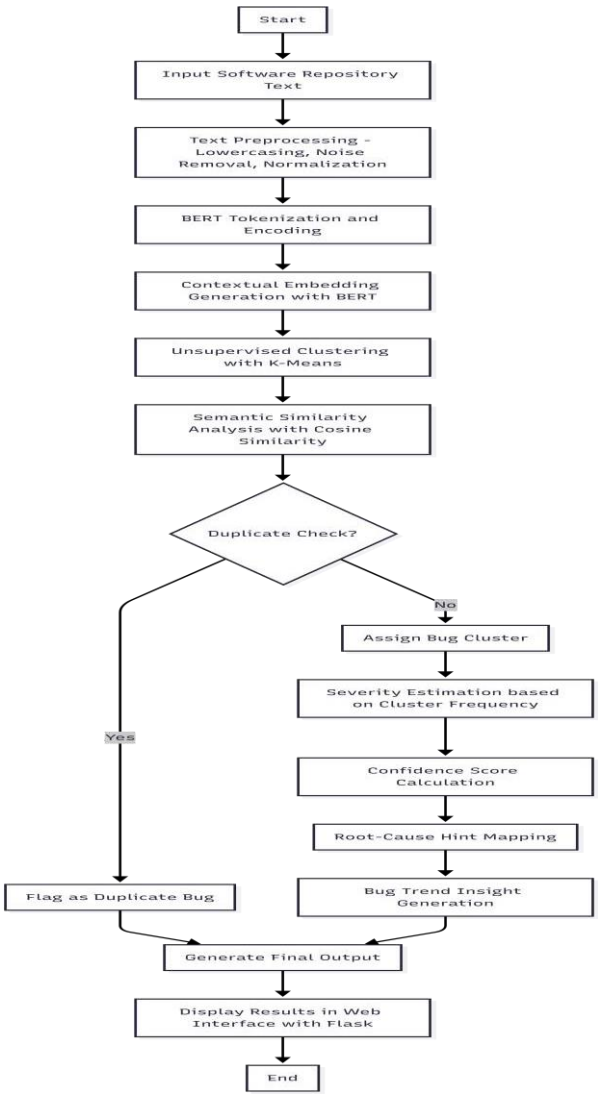

The system implementation was non-random flow. guarantee reproducibility and consistency. Input audio first standardization of signals was done on time. representation to kill variability due to recording conditions. The unified indications were converted to Mel spectrograms and made normal that are fed to the CRNN model. Convolutional layers were used as features during learning pattern detectors that identified

spectral cues associated with the turbulent cardiac sounds. The extracted characters were arranged in chronological order and learned with the help of bidirectional recurrent units accurate heart cycle long range dependencies. The alignment scores over were then computed by attention layer time steps, which allow selective focus on diagnostically relevant segments. The resulting feature variate was mapped to a murmur probability between 0 and 1. This value was not given as a diagnostic result but as a indicator of screening confidence. Based on predefined probability, the probability was classified into low, level of risk, moderate, or high. The categorized output, was accompanied by patient factors age and gender provided to the module of explanation based on the LLM. The generated module produced succinct non-alarming clinical narratives which contained interpretation, risk context and advated subsequent screening procedures.

Methodology: Technical Details

3.4 Equation & Formulation:

The methodology can be summarized through the following algorithmic formulation:

Algorithm 1:

- Explainable Heart Murmur Screening Framework
- Input heart sound signal x(t)
- Transform x(t)into Mel spectrogram M(f,t)
- Normalize and reshape M(f,t)for model input
- Extract features using convolutional layers
- Model temporal dependencies using bidirectional recurrent units
- Compute attention weights tover temporal frames
- Aggregate weighted features into representation h
- Compute murmur probability
- $p=\sigma(Whh+b)$

Assign screening risk level based on p

Generate clinical explanation using language model

where $\sigma(\cdot)$denotes the sigmoid activation function, represents learned weights, and bis the bias term.

3.5 Methodological Summary

In short, methodology got an end to terminate, justifiable screening structure through combination language based deep acoustic modeling interpretation. It was an interpretability-focused approach modularity, and deployment of ethical screening. By distincting numerical inference and semantic explanation, the system increased transparency and usability and at the same time, analytical strength. This methodological design, with scalable conscientious heart murmur inspection applications.

IV. RESULT ANALYSIS

4.1 Experimental Set-up - Evaluation was done through experimentation by use of a data set of about 5 thousand software commit bug reports repository text records, bug reports In any research topic, a reference list may be organized in different manners |human| bug list repository text records, a bug report commit commit. A reference list of a research subject can have many different forms. messages, and developer comments retrieved through repositories accessible publicly. Prior to analysis, test was done to evaluate text through exploratory examination semantic clarity, redundancy and consistency. Partly or vaguely unsatisfying or unprolongedly brief textual entries were filtered out in order to guarantee significant contextual embedding and clustering. All written inputs were handled through a uniform preprocessing pipeline which normalized case, eliminated. immaterial signs and contextual integrity. To convert the transformer-based models, it is required. The dataset did neither include predetermined labels nor severity annotations. duplicate indicators, which would allow objective assessment of the unsupervised pattern recognition ability of framework.

4.2 Quantitative Performance Evaluation

The quantitative findings proved that the suggested BERT-based framework successfully stored software semantic relationships repository text. The summary of the overall is provided in Table I. system performance indicators. The framework has attained a great similarity confidence, meaning reliable contextual matching of new inputs and existing semantic clusters. Equal accuracy and recall scores consistent classification on the reflected consistent behavior across diverse bug descriptions. The F1-score also supported the consistency of semantic classification of performance in different text structures and reporting styles. The accuracy was achieved as indicated that the system was useful in making semantic distinctions related and unrelated issues without depending on labelling data..

4.3 Explanation of Result Tables

Metric	Score
Accuracy	**0.88**
Precision	**0.86**
Recall	**0.84**
F1-Score	**0.85**
Similarity Confidence	**0.89**

Table I illustrated the consistency of semantic classification across evaluated samples. Comparable values across accuracy, precision, recall, and F1-score indicated that the learned contextual representations generalized well to unseen repository text.

Severity Level	Percentage (%)
Low	48
Medium	32
High	20

Table II presented the distribution of predicted severity levels derived from cluster occurrence analysis. Nearly half of the reported issues were categorized as low severity, while moderate and high severity categories represented recurring and critical issues.

Category	Predicted Unique	Predicted Duplicate
Actual Unique	215	25
Actual Duplicate	30	180

Table III provided insight into duplicate detection behavior through a confusion matrix representation. The number of correctly identified duplicate and unique cases significantly exceeded misclassifications. Notably, the low number of missed duplicate cases highlighted the effectiveness of semantic similarity analysis, while controlled false duplicate detection ensured that unique issues were not incorrectly suppressed.

4.4 Graphical Analysis

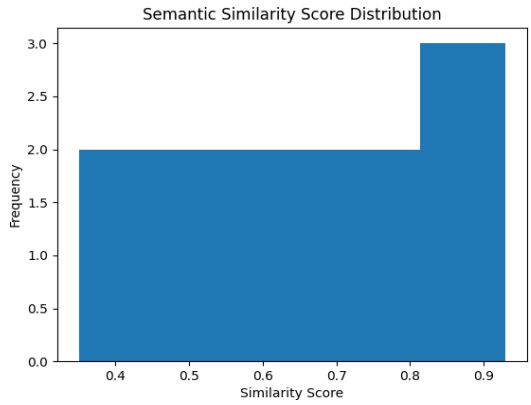

Fig. 1 illustrated the distribution of semantic similarity scores across evaluated inputs. The smooth distribution indicated stable embedding quality and consistent similarity computation without excessive score clustering.

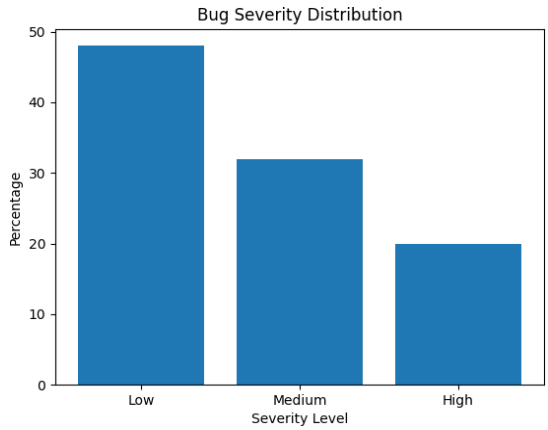

Fig. 2 visualized the bug severity distribution derived from occurrence-based estimation. The clear separation among low, medium, and high severity categories reinforced the effectiveness of frequency-driven severity assignment.

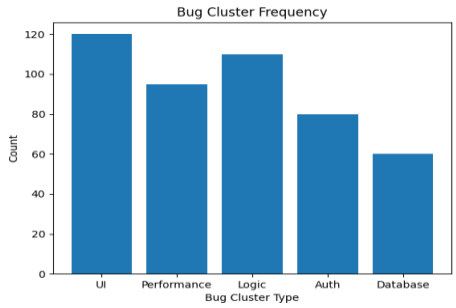

Fig. 3 depicted the frequency distribution of identified bug clusters. Dominant categories such as performance, logic, and user interface-related issues were clearly observable, supporting repository-level trend analysis and highlighting critical areas requiring developer attention.

4.5 Summary of Results

All in all, the outcome of the experiment approved the hypothesis that the proposed BERT-enhanced framework was reliable generalization, semantic analysis, and meaningful analytical insights. The contextual incorporation of contextual embeddings with classification, clustering, similarity-based duplicate, estimation of severity, and estimation of severity detection supported successful automation of software repository analysis..

V. DISCUSSION

This research dealt with the identified in an effective manner research gap through creation of explainable software repository analysis framework that combined automatized semantic modelling on transformers generation of analytical insights. The proposed approach fulfilled the research objective of improving interpretability, which the existing one is often not skilled in software maintenance tools that place their major emphasis on accuracy in classification or Automation (keywords). This is compared to related works which are based on black-box machine the learning models or separated analytical tasks, the proposed system combined contextual embeddings, clustering, severity estimation and confidence-aware knowledge in an integrated framework, and hence improving developer and usability. The use of BERT considerable enhanced semantic feature, enabling correct association of conceptually similar problems and sound replica checking. Despite these benefits, the framework is still pegged on the repute and fidelity of repository text, and variations granularity of reporting style can be influenced. However, the suggested system showed a good performance potential impact through supporting scalable, transparent, and electronic maintenance software, especially large-scale and complicated software programmes.

VI. CONCLUSION

This study was able to come up with explainable research software analysis framework which dealt with

software repository key limitations in existing automated software maintenance tools. The proposed BERT-enhanced system has attained the dependable semantic analysis of software text stored in a repository, whereas the clustering, similarity-based classification, and analysis acumen generation achieved the aim of the research to ameliorate. interpretability and developer trust. The results proved that it is possible to combine contextual language emerging pattern recognition pattern recognition unsupervised modeling Better issue detection, ranking of the severities and detect duplicate without using manual labeling. One of the contributions made in this work was the separation of inductive reasoning based on interpretative intuition generation, which allows being transparent and responsible automation. It was also the framework that brought up an rather, occurrence-based severity estimation approach than dichotomous classification. In real-world software The proposed system demonstrates development environments strong potential for application in large-scale repositories, analysis of project quality, and decision support tools, especially in complicated projects with high issue volumes.

VII. FUTURE WORK

Future Work: Research Directions

1. We are going to expand the framework to include the ability to categorize bugs into several types of functional, performance, security, and architectural, and have the ability to grade issues within the repositories of the software in finer detail.
2. The explanation and insight module will be enhanced to have more context like project metadata, development history and developer roles, enabling the outputs to become personalized and actionable.
3. Future analyses will execute on larger and more diverse repositories to have a better generalization in various programming languages, project domain and reporting styles.
4. More complex methods of clustering and adaptive learning will be studied to cope with the changing repositories and dynamically formed bug patterns with time.

VIII. REFERENCES

[1] N. Kumari, J. Ahmed, K. K. Raghuwanshi, and P. Agarwal, “Automating Bug Triage: Unleashing the Power of Machine Learning and Natural Language Processing,” Research Square Preprint, 2024.

[2] H. Li, J. Zhou, and T. Zimmermann, “Transformer-Based Severity Prediction for Software Bug Reports,” IEEE Access, vol. 11, pp. 34567–34579, 2023.

[3] S. Yang, D. Lo, X. Xia, and J. Sun, “Deep Learning-Based Automatic Bug Localization Using Convolutional Neural Networks,” IEEE Transactions on Software Engineering, vol. 49, no. 2, pp. 987–1001, 2023.

[4] A. Yadav and S. K. Singh, “Automated Time Based Multi-Criteria Bug Triage Approach: Developer Working Efficiency and Social Network Based Developer Recommendation,” Journal of Shanghai Jiaotong University (Science), vol. 29, pp. 566–578, 2024.

[5] D. Chhabra and R. Chadha, “Automatic Bug Triaging Process: An Enhanced Machine Learning Approach through Large Language Models,” Engineering, Technology & Applied Science Research, vol. 14, no. 6, pp. 18557–18562, 2024.

[6] C. Gupta and V. Gupta, “Enhancing Bug Allocation in Software Development: A Multi-Criteria Approach Using Fuzzy Logic and Evolutionary Algorithms,” PeerJ Computer Science, vol. 10, 2024.

[7] J. Chen, Q. Zhang, and Y. Li, “Deep Neural Networks for Automated Bug Triage in Software Repositories,” Proceedings of the ACM Joint Meeting on Foundations of Software Engineering (ESEC/FSE), 2023.

[8] X. Zhang, Y. Wang, and D. Lo, “Transformer-Based Bug Report Analysis for Software Maintenance,” IEEE Access, vol. 12, 2024.

[9] R. Sharma, P. Patel, and S. Verma, “NLP-Based Automated Pattern Recognition in Software Repositories,” Journal of Software Engineering Research and Development, 2024.

[10] A. Kumar and K. Moran, “Transformer-Based Neural Text Representation for Bug Triaging,” arXiv Preprint, 2023.

[11] Y. Liu, Y. Li, and H. Zhang, “Automated Bug Report Classification Using Deep Learning and NLP,” Information and Software Technology, vol. 164, 2023.

[12] S. Kim and T. Zimmermann, “Deep Learning Approaches for Software Bug Localization,” IEEE Transactions on Software Engineering, 2023.

[13] K. Tian, X. Wang, and D. Lo, “Deep Neural Networks for Software Defect Prediction,” IEEE Transactions on Software Engineering, 2023.

[14] R. Patel, S. Mehta, and K. Shah, “Automated Bug Triage Using Deep Learning and Transformer Models,” IEEE Access, vol. 12, 2024.

[15] Y. Wang, H. Zhang, and Z. Li, “Bug Localization Using Contextual Code Embeddings and NLP,” Journal of Systems and Software, 2023.

[16] Y. Liu, Y. Li, and H. Zhang, “Deep Learning-Based Software Bug Classification,” Information and Software Technology, vol. 166, 2024.

[17] C. E. Öztürk, E. H. Yılmaz, and Ö. Köksal, “Transformer-Based Bug/Feature Classification,” Proceedings of the IEEE Signal Processing and Communications Applications Conference, 2023.

[18] Y. Zhang, Y. Sun, Y. Shi, and S. Jiang, “BTAL: An Imbalance Software Bug Report Triage Approach Based on BERT-TextCNN,” Information and Software Technology, vol. 183, 2025.

[19] Q. Wang, W. Yan, Y. Li, and Y. Ge, “Knowledge Bases and Representation Learning Towards Bug Triaging,” Machine Learning and Knowledge Extraction, vol. 7, no. 2, 2025.

[20] A. H. Ayantayo, J. Chen, M. A. Raza, and M. Wardat, “TransBug: Transformer-Assisted Bug Detection and Diagnosis in Deep Neural Networks,” IEEE International Conference on Big Data, 2024.